



北京大学高性能计算与应用中心

Center for Energy-efficient Computing and Applications

Introduction to Parallel & Distributed Computing Programming Models

Lecture 2, Spring 2014


Instructor: 罗国杰

gluo@pku.edu.cn

Example: Compute π


◆ Computing π

$$\left. \begin{array}{l} \arctan(1) = \pi/4 \\ \arctan(0) = 0 \\ \frac{d}{dx} \arctan(x) = 1/(1+x^2) \end{array} \right\} \Rightarrow \pi = \int_0^1 \frac{4}{1+x^2} dx$$



Example: Sequential Code

```
double compute_pi(int n) {  
    double sum = 0.0;  
    for (int i = 0; i < n; i++) {  
        double x = (i + 0.5) / n;  
        sum += 1.0 / (1.0 + x*x);  
    }  
    double pi = 4.0 * sum / n;  
    return pi;  
}
```



POSIX Thread Version (1/2)

// global variables

int thread_count;

int n;

double* local_sum;

// multithreaded version

double compute_pi () {

 thread_handles = (pthread_t*) malloc (thread_count*sizeof(pthread_t));

 local_sum = (double*) malloc (thread_count*sizeof(double));

 // parallel part

 for (thread = 0; thread < thread_count; thread++)

 pthread_create(&thread_handles[thread], NULL, thread_sum, (void*)thread);

 for (thread = 0; thread < thread_count; thread++)

 pthread_join(thread_handles[thread], NULL);

 // sequential part

 double sum = 0.0;

 for (thread = 0; thread < thread_count; thread++)

 sum += local_sum[thread];

 double pi = 4.0 * sum / n;

 return pi;

}

POSIX Thread Version (2/2)

```
void* thread_sum(void* rank) {  
    int my_rank = (int) rank;  
    double my_sum = 0.0;  
    // domain decomposition  
    for (int i = my_rank; i < n; i += thread_count) {  
        double x = (i + 0.5) / n;  
        my_sum += 1.0 / (1.0 + x * x);  
    }  
    local_sum[my_rank] = my_sum;  
    return NULL;  
}
```

OpenMP Version

```
/* compile as $> gcc -fopenmp -lm
 * run as      $> OMP_NUM_THREADS=2 ./a.out
 */
double compute_pi(int n) {
    int i;
    double sum = 0.0;
    #pragma omp parallel for reduction(+: sum) schedule(static)
    for (i = 0; i < n; ++i) {
        double x = (i + 0.5) / n;
        sum += 1.0 / (1.0 + x * x);
    }
    double pi = 4.0 * sum / n;
    return pi;
}
```

MPI Version (1/2)

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
int main( int argc, char *argv[] ) {
    int n, my_rank, numprocs, i;
    double mypi, pi, h, sum, x;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&my_rank);
    while (1) {
        if (my_rank == 0) {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d",&n);
        }
        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
```

MPI Version (2/2)

```
if (n == 0) break;
else {
    sum = 0.0;
    for (i = my_rank; i <= numprocs) {
        x = (i + 0.5) / n;
        sum += 4.0 / (1.0 + x*x);
    }
    mypi = sum / n;
    MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
               MPI_COMM_WORLD);
    if (my_rank == 0) printf("pi is approximately %.16f\n", pi));
}
MPI_Finalize();
return 0;
}
```


OpenCL Version

```
/* to be executed on device */
__kernel void pi(const int  niters,  const float  step_size,
                __local  float*  local_sums,  __global float*  partial_sums) {
    int num_wrk_items  = get_local_size(0);
    int local_id        = get_local_id(0);
    int group_id        = get_group_id(0);

    float x, accum = 0.0f;
    int i, istart, iend;

    istart = (group_id * num_wrk_items + local_id) * niters;
    iend   = istart+niters;

    for(i= istart; i<iend; i++) {
        x = (i+0.5f)*step_size;
        accum += 4.0f/(1.0f+x*x);
    }

    local_sums[local_id] = accum;
    barrier(CLK_LOCAL_MEM_FENCE);

    reduce(local_sums, partial_sums);
}

/* to be executed on host */
pi_res = 0.0f;
for (unsigned int i = 0; i < nwork_groups; i++) { pi_res += h_psum[i]; }
pi_res *= step_size;
```

Outline

- ◆ **Parallel programming models and architectures**
- ◆ **Memory limitations**

Architectural Trends

◆ Greatest trend in VLSI is an increase in the exploited parallelism

- Up to 1985: **bit-level parallelism**: 4-bit -> 8-bit -> 16-bit
 - slows after 32 bits
 - adoption of 64-bit now under way
- Mid 80s to mid 90s: **instruction-level parallelism**
 - pipelining and simple instruction sets (RISC)
 - on-chip caches and functional units => superscalar execution
 - greater sophistication: out of order execution
- Nowadays:
 - hyper-threading
 - multi-core

Evolution of Architectural Models

- ◆ Historically (1970s – early 1990s), each parallel machine was unique, along with its programming model and language
 - Architecture = prog. model + comm. abstraction + machine organization
- ◆ Throw away software & start over with each new kind of machine
 - Dead Supercomputer Society
 - <http://www.paralogos.com/DeadSuper/>
- ◆ Nowadays we separate the **programming model** from the underlying **parallel machine architecture**
 - 3 or 4 dominant programming models
 - Dominant: shared address space, message passing, data parallel
 - Others: data flow, systolic arrays

Programming Model for Various Architectures

- ◆ **Programming models specify communication and synchronization**
 - **Multiprogramming: no communication/synchronization**
 - **Shared address space: like bulletin board**
 - **Message passing: like phone calls**
 - **Data parallel: more regimented, global actions on data**
- ◆ **Communication abstraction: primitives for implementing the model**
 - **Play the role like the instruction set in a uniprocessor computer**
 - **Supported by HW, by OS, or by user-level software**
- ◆ **Programming models are the abstraction presented to programmers**
 - **Write portably **correct** code that runs on many machines**
 - **Write **fast** code requires tuning for the architecture**
 - **Not always worthy of it – sometimes programmer time is more precious**

Aspects of a Parallel Programming Model

◆ Control

- How is **parallelism** created?
- In what **order** should operations take place?
- How are different threads of control **synchronized**?

◆ Naming

- What data is private vs. shared?
- How is shared data accessed?

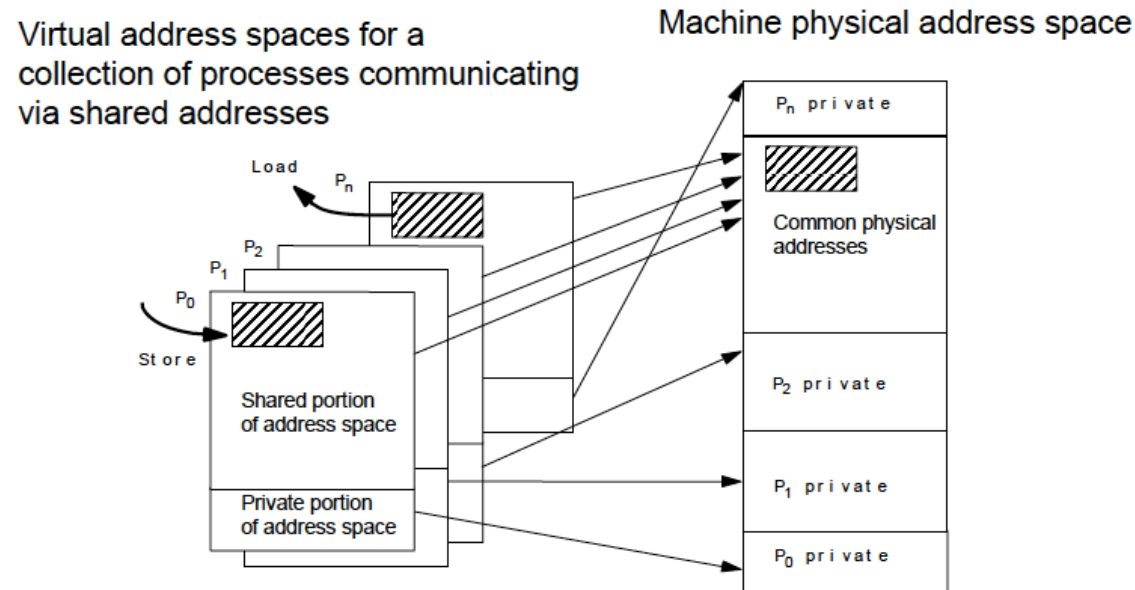
◆ Operations

- What operations are atomic?

◆ Cost

- How do we account for the cost of operations?

Programming Models: Shared Address Space (SAS)



◆ Programming model

- Process: virtual address space plus one or more threads of control
- Portions of address spaces of processes are shared
- Writes to shared address visible to all threads (in other processes as well)

◆ Natural extension of uniprocess model

- conventional memory operations for communication
- special atomic operations for synchronization

SAS Machine Architecture (1/2)

◆ Motivation: programming convenience

■ Location transparency

- Any processor can directly reference any memory location
- Communication occurs implicitly as result of loads and stores

■ Extended from time-sharing on uni-processors

- Processes can run on different processors
- Improved throughput on multi-programmed workloads


◆ Communication hardware also natural extension of uniprocessor

- Addition of processors similar to memory modules, I/O controllers

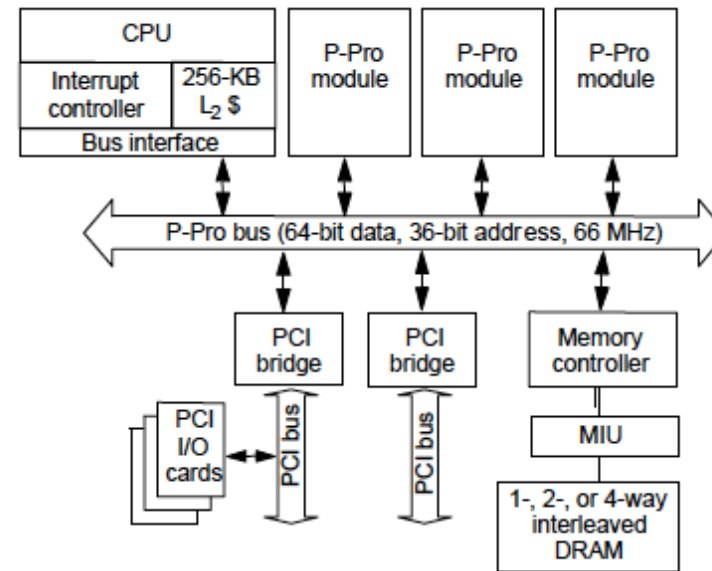

SAS Machine Architecture (2/2)

◆ One representative architecture: SMP

- Used to mean *Symmetric MultiProcessor*
 - All CPUs had equal capabilities in every area, e.g., in terms of I/O as well as memory access
- Evolved to mean *Shared Memory Processor*
 - Non-message-passing machines (included crossbar as well as bus based systems)
- Now it tends to refer to *bus-based shared memory machines*
 - Small scale: < 32 processors typically




Example: Intel Pentium Pro Quad



- All coherence and multiprocessing glue in processor module
- Highly integrated, targeted at high volume
- Low latency and high bandwidth


Scaling Up: More SAS Machine Architectures (1/2)



◆ Dance-hall

- Problem: interconnect cost (crossbar) or bandwidth (bus)
- Solution: scalable interconnect network -> bandwidth scalable
- Latencies to memory uniform, but uniformly large (**Uniform Memory Access, UMA**)
- Caching is key: coherence problem

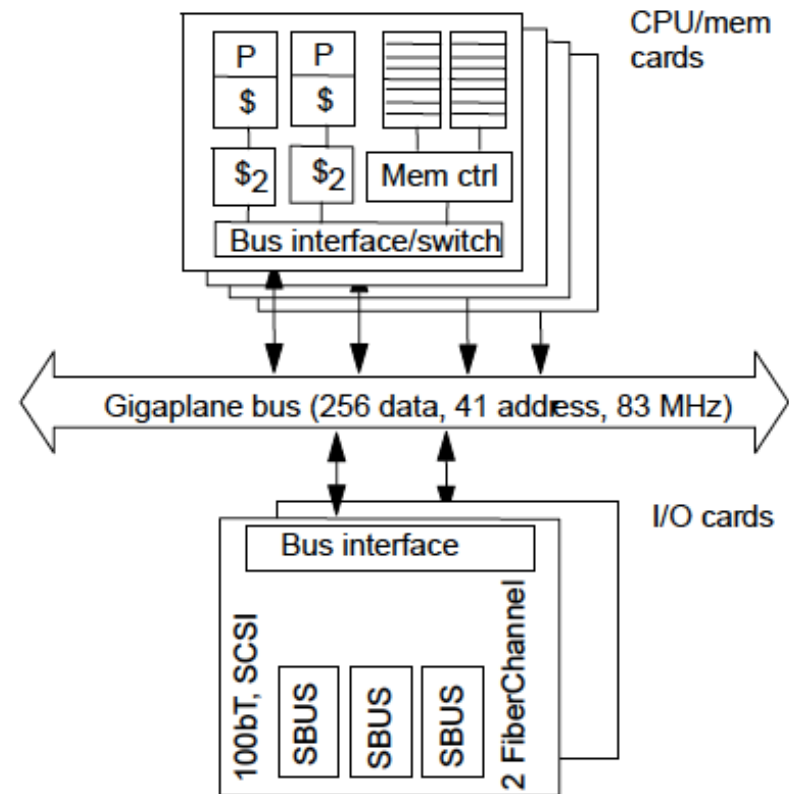
Scaling Up: More SAS Machine Architectures (2/2)



◆ Distributed shared memory (DSM) or non-uniform memory access (NUMA)


- Non-uniform time for the access to data in local memory and remote memory
- Caching of non-local data is key
 - coherence cost

Example: SUN Enterprise




- 16 cards of either type: processors + memory, or I/O
- All memory accessed over bus, so symmetric
- Higher bandwidth, higher latency bus

Example: Cray T3E



- Scale up to 1024 processors, 480MB/s links
- Memory controller generates comm. request for nonlocal references
- No hardware mechanism for coherence (SGI Origin etc. provide this)

Programming Models: Message Passing



◆ Programming model


- Directly access only private address space (local memory), communicate via explicit messages
 - Send specifies data in a buffer to transmit to the receiving process
 - Recv specifies sending process and buffer to receive data
- In the simplest form, the send/rcv match achieves pair-wise synchronization

◆ Model is separated from basic hardware operations


- Library or OS support for copying, buffer management, protection
- Potential high overhead: large messages to amortize the cost

Message Passing Architectures

- ◆ Complete processing node (computer) as building block, including I/O
 - Communication via explicit I/O operations
 - Processor/Memory/IO form a processing node that cannot directly access another processor's memory
- ◆ Each “node” has a network interface (NI) for communication and synchronization



DSM vs. Message Passing

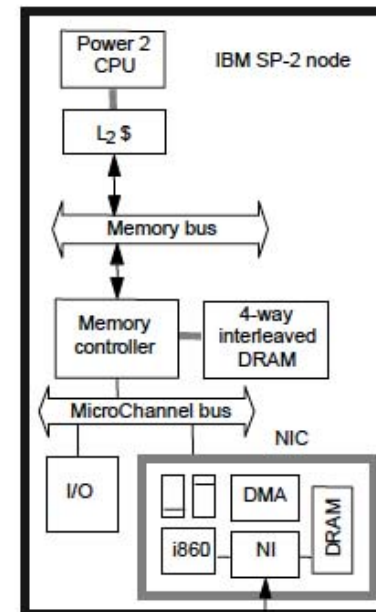



- ◆ **High-level block diagrams are similar**
- ◆ **Virtualization is possible, with some performance penalty, i.e.**
 - MP programs can run on DSM architectures
 - DSM programs can run on MP architectures, with proper library or OS support for memory sharing
- ◆ **Implications of DSM and MP on architectures**
 - **Fine-grained hardware supports for DSM**
 - Communication integrated at I/O level for MP, needn't be into memory system
 - MP can be implemented as middleware (library)
 - MP has better scalability
 - ◆ MP machines are easier to build than scalable address space machines

Example of MP: IBM SP-2




General interconnect network formed from 8-port switches

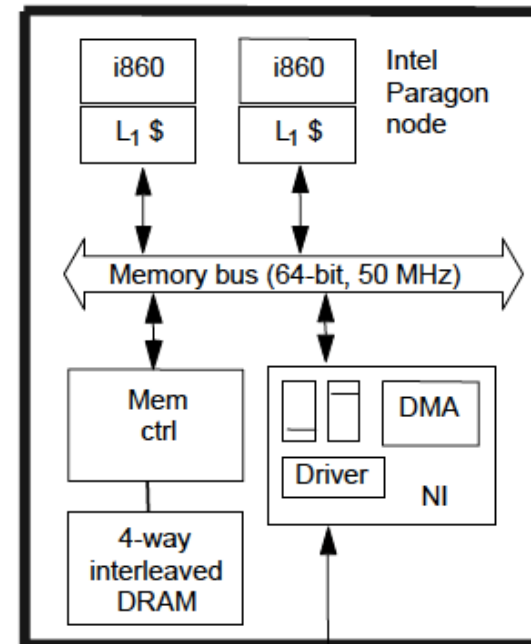


- Each node is a essentially complete RS6000 workstation;
- Network interface integrated in I/O bus (bw limited by I/O bus).


Example: Intel Paragon



Sandia's Intel Paragon XP/S-based Supercomputer



2D grid network
with each processing node
attached to a switch



Toward Architectural Convergence

- ◆ **Convergence in hardware organizations**
 - Tighter NI integration for MP
 - Hardware SAS passes messages at lower level
 - Cluster of workstations/SMP become the most popular parallel architecture for parallel systems
- ◆ **Programming models distinct, but organizations converging**
 - Nodes connected by general network and communication assists
 - Implementations also converging, at least in high-end machines

Programming Model: Data Parallel (1/2)

- ◆ Operations performed in parallel on each element of data structure
- ◆ Logically single thread of control (sequential program)
- ◆ Conceptually, a processor associated with each data element
- ◆ Coordination is implicit – statements executed synchronously
- ◆ Example:

```
float x[100];
```


```
for (i=0; i<100; i++)  
    x[i] = x[i] + 1;
```

➔ $x = x + 1;$

Programming Model: Data Parallel (2/2)

◆ Architecture model

- A control processor issues instructions
- Array of many simple cheap processors – processing element (PE) – each with little memory
- A interconnect network that broadcasts data to PEs, communication among PEs, and cheap synchronization



◆ Motivation

- Give up flexibility (different instructions in different processors) to allow a much larger number of processors
- Target at limited scope of applications


◆ Applications

- Finite differences, linear algebra
- Document searching, graphics, image processing

A Case of DP: Vector Machine

◆ Vector machine

- Multiple functional units
- All performing the same operation
- Instructions may be of very high parallelism (e.g. 64-way) but hardware executes only a subset in parallel at a time



◆ Historically important, but overtaken by MPPs (massive parallel processors) in the 90s

◆ Re-emerging in recent years

- At a large scale in the Earth Simulator (NEC SX6) and Cray X1
- At a small scale in SIMD media extensions to microprocessors
 - SSE (Streaming SIMD Extensions), SSE2 (Intel: Pentium/IA64)
 - AltiVec (IBM/Motorola/Apple: PowerPC)
 - VIS (Sun: Sparc)

◆ Enabling technique

- Compiler does some of the difficult work of finding parallelism

Flynn's Taxonomy


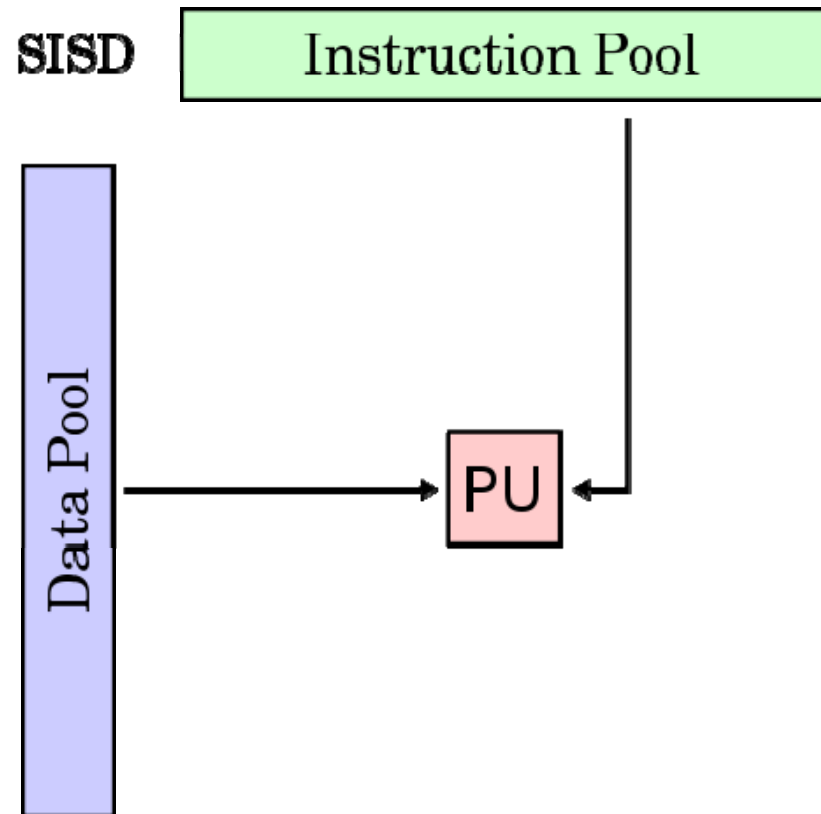
- ◆ **A classification of computer architectures based on the number of streams of instructions and data:**

		Instruction stream	
		single	multiple
Data stream	single	SISD Single Instruction, Single Data	MISD Multiple Instructions, Single Data
	multiple	SIMD Single Instruction, Multiple Data	MIMD Multiple Instructions, Multiple Data

- ◆ **Program model converges with SPMD (single program multiple data)**

SISD Architecture

Example: single-core computers

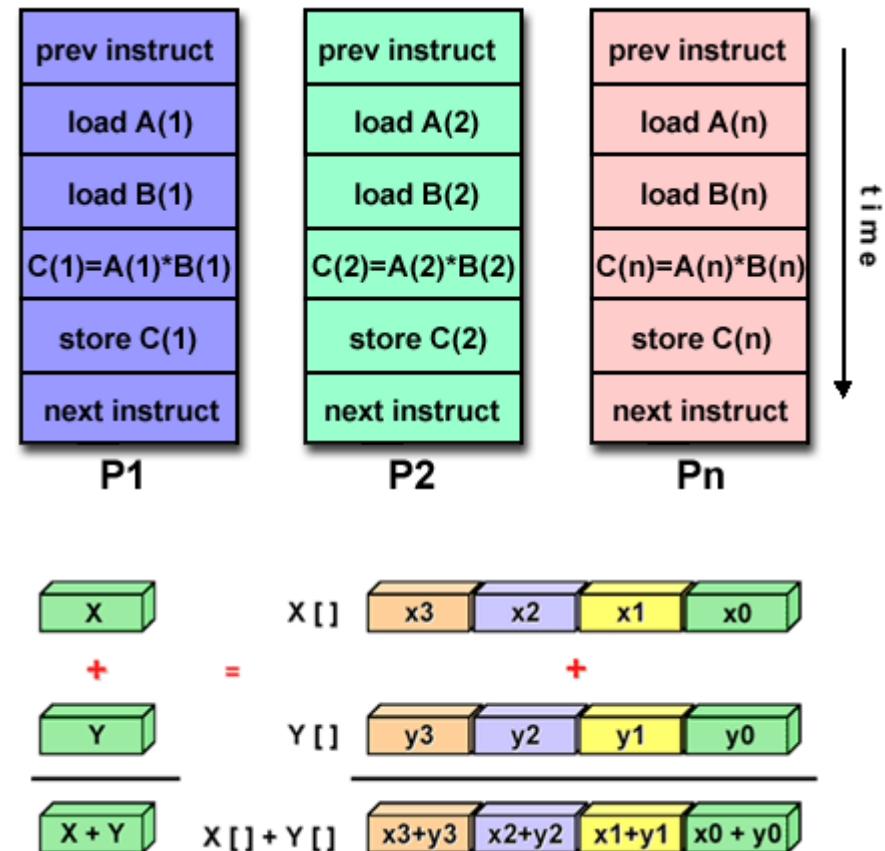



Source: http://en.wikipedia.org/wiki/Flynn's_taxonomy

Source: Blaise Barney (LLNL), Introduction to Parallel Computing

SIMD Architecture

Example: vector processors, GPUs

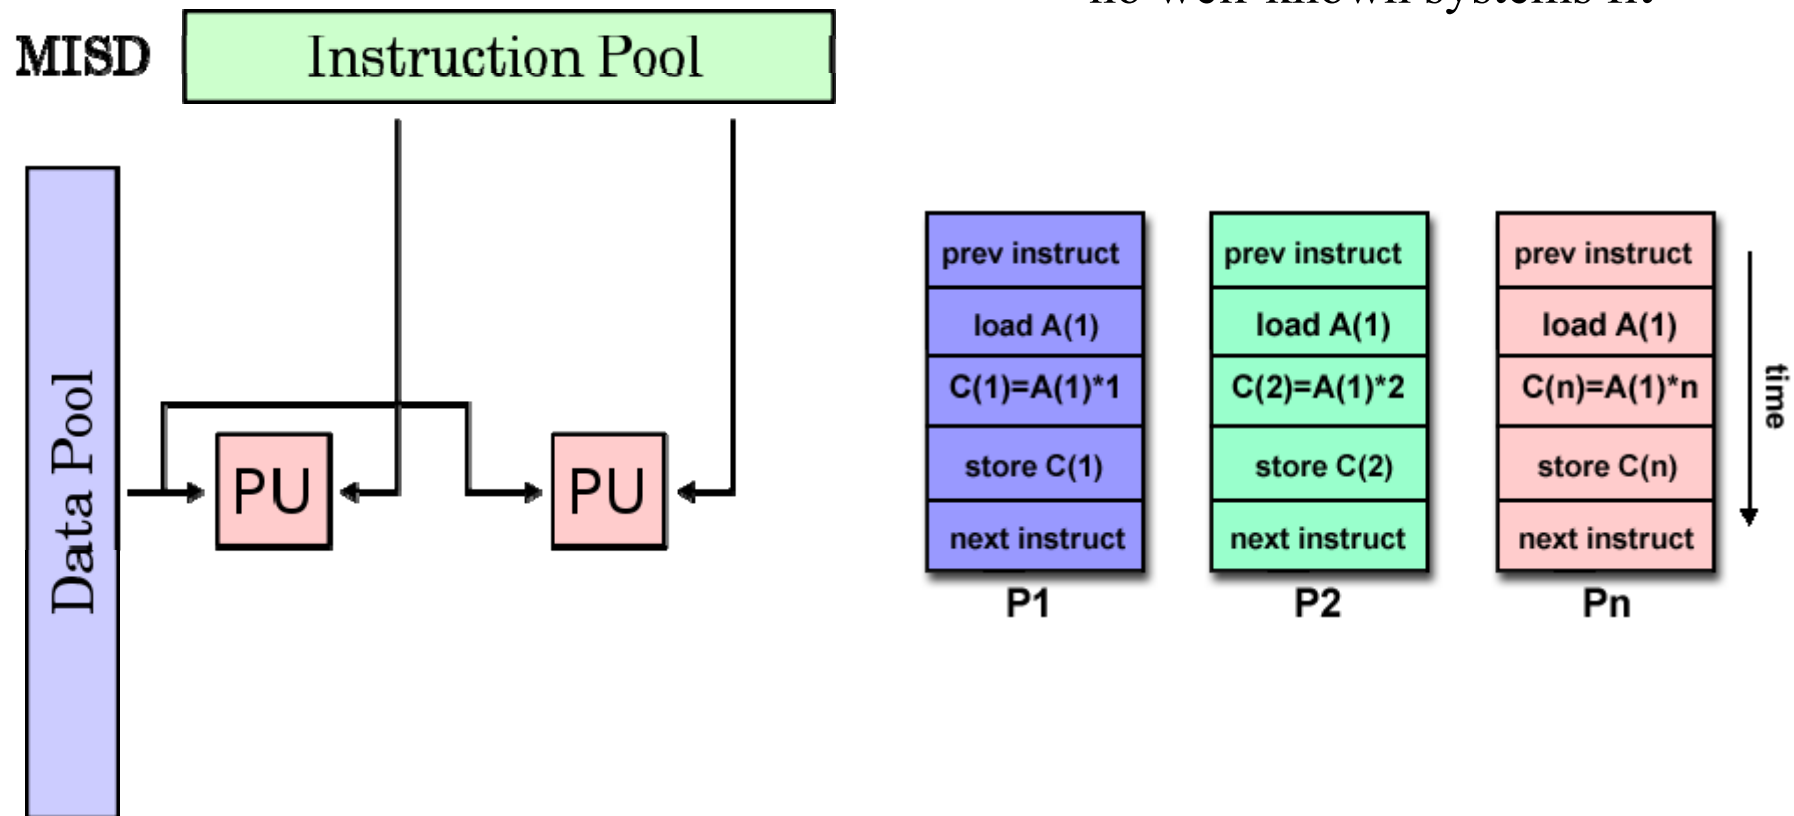


Source: http://en.wikipedia.org/wiki/Flynn's_taxonomy

Source: Blaise Barney (LLNL), Introduction to Parallel Computing

MISD Architecture


* no well-known systems fit



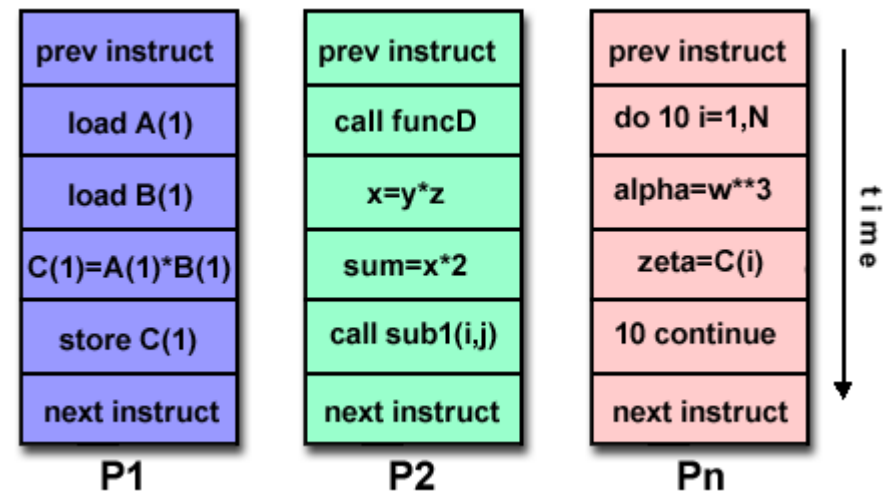
Source: http://en.wikipedia.org/wiki/Flynn's_taxonomy

Source: Blaise Barney (LLNL), Introduction to Parallel Computing

MIMD Architecture



Example: modern parallel systems



Source: http://en.wikipedia.org/wiki/Flynn's_taxonomy

Source: Blaise Barney (LLNL), Introduction to Parallel Computing

A Further Breakdown of MIMD

◆ Shared-memory architecture

- **Symmetric multiprocessors (SMP), or uniform memory access (UMA)**
- **Nonuniform memory access (NUMA) systems**
 - **Cache-coherent NUMA (ccNUMA)**

◆ Distributed-memory architecture

- **Massively parallel processors (MPP)**
 - **Tightly coupled, specialized processors and network infrastructure**
- **Clusters**
 - **Off-the-shelf computers connected by an off-the-shelf network**

◆ Hybrid distributed-shared memory architecture


◆ Grids

- **Distributed, heterogeneous resources connected by LANs and/or WANs**


From T.G. Mattson, et al., “Patterns for Parallel Programming,” Addison-Wesley Professional, 2004.

Shared-Memory Architecture

Symmetric Multiprocessors (SMP)




Nonuniform Memory Access (NUMA)



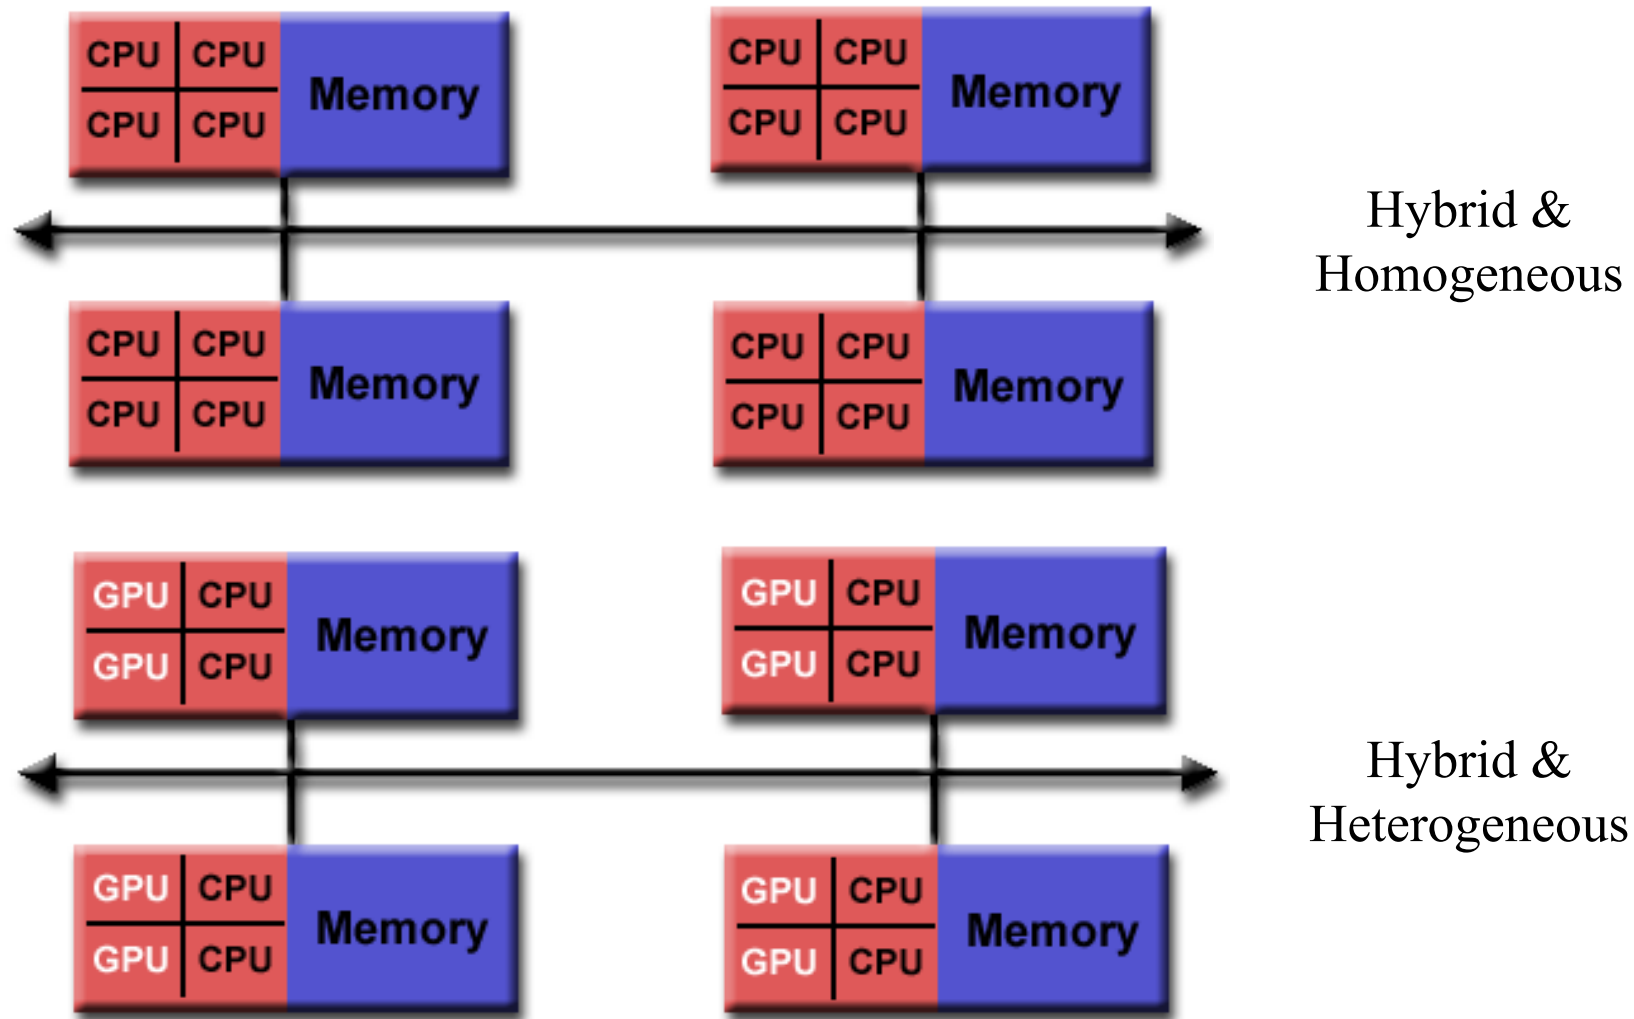
Source: Blaise Barney (LLNL), Introduction to Parallel Computing

Distributed-Memory Architecture



Source: Blaise Barney (LLNL), Introduction to Parallel Computing

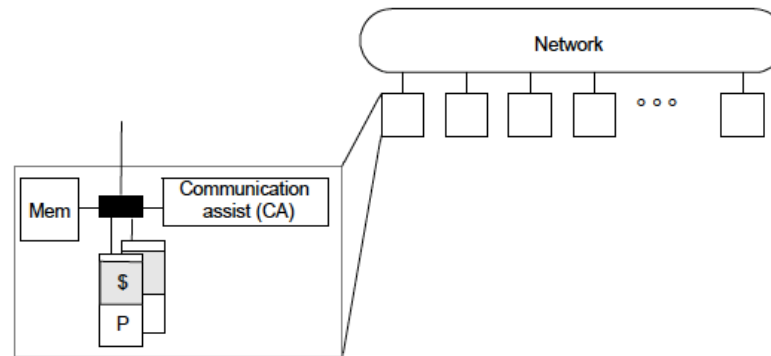
Hybrid Architecture



Source: Blaise Barney (LLNL), Introduction to Parallel Computing

Convergence: Generic Parallel Architecture

◆ A generic modern multiprocessor



◆ Node: processor(s), memory, plus communication assist

- Network interface and communication controller

◆ Scalable network

◆ => Convergence allows lots of innovation, now within the same framework

- integration of assist within node, what operation, how efficiently...

Outline

- ◆ **Parallel programming models and architectures**
- ◆ **Memory limitations**

Limitations of Memory System Performance

- ◆ Memory system, and not processor speed, is often the bottleneck for many applications.
- ◆ Memory system performance is largely captured by two parameters, latency and bandwidth.
- ◆ **Latency** is the time from the issue of a memory request to the time the data is available at the processor.
- ◆ **Bandwidth** is the rate at which data can be pumped to the processor by the memory system.

Difference between Bandwidth and Latency

◆ An example of a fire-hose


- If the water comes out of the hose two seconds after the hydrant is turned on, latency = two seconds.
- Once the water starts flowing, if the hydrant delivers water at the rate of 20 liters/second, bandwidth = 5 liters/second.

◆ Serve different purposes

- If you want immediate response from the hydrant, it is important to reduce latency.
- If you want to fight big fires, you want high bandwidth.

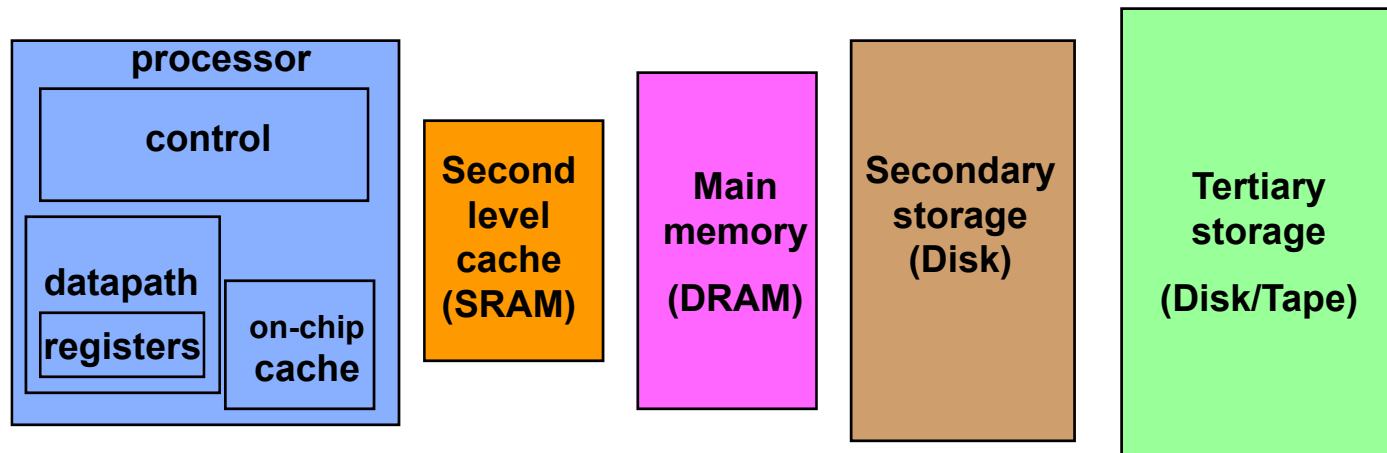
Processor-DRAM Gap (Latency)

- ◆ Memory hierarchies are getting deeper
 - Processors get faster more quickly than memory



Memory Hierarchy

- ◆ Most programs have a high degree of **locality** in their accesses
 - spatial locality: accessing things nearby previous accesses
 - temporal locality: reusing an item that was previously accessed
- ◆ Memory hierarchy tries to exploit locality to improve average



Speed	1ns	10ns	100ns	10ms	10sec
Size	KB	MB	GB	TB	PB

Memory Latency: An Example

- ◆ **A processor operating at 1 GHz (1 ns clock) connected to a DRAM with a latency of 100 ns (no caches).**
- ◆ **Assume the processor has 2 multiply-add units; capable of executing 4 instructions in each clock cycle**
 - **The peak processor rating is 4 GFLOPS.**
 - **Every time a memory request is made, the processor must wait 100 cycles before it can process the data.**

Memory Latency: An Example (Cont'd)

- ◆ **On the previous architecture, consider the problem of computing a dot-product of two vectors.**
 - **Performs one multiply-add on each pair of vector elements**
 - Each floating point operation requires one data fetch.
 - **Limited to one floating point operation every 100 ns, or a speed of 10 MFLOPS,**
 - A very small fraction of the peak processor rating!

Improving Effective Memory Latency Using Caches

- ◆ **Caches are small and fast memory elements between the processor and DRAM.**
 - Acts as a low-latency high-bandwidth storage.
 - If a piece of data is repeatedly used, the effective latency of this memory system can be reduced by the cache.
- ◆ **The fraction of data references satisfied by the cache is called the cache *hit ratio*.**
- ◆ **Cache hit ratio achieved by a code on a memory system often determines its performance.**

Impact of Caches: Example

- ◆ Introduce a cache of size 32 KB with a latency of 1 ns or one cycle in previous example.
- ◆ Assume we multiply two matrices A and B of dimensions 32×32 .
 - Both A and B (1K each), as well as the result matrix C can be in cache
- ◆ Fetching the two matrices into the cache corresponds to fetching 2K words, which takes approximately 200 μ s.
- ◆ Multiplying two $n \times n$ matrices takes $2n^3$ operations.
 - \Rightarrow 64K operations, which can be performed in 16K cycles (or 16 μ s) at 4 instructions per cycle.
- ◆ The total time for the computation = $200 + 16 \mu$ s.
 - A peak computation rate of 64K/216 or 303 MFLOPS
 - Much better!

Impact of Caches

- ◆ Repeated references to the same data item correspond to temporal locality.
- ◆ In our example, we had $O(n^2)$ data accesses and $O(n^3)$ computation. This asymptotic difference makes the above example particularly desirable for caches.
- Data reuse is critical for cache performance.

Impact of Memory Bandwidth

- ◆ **Memory bandwidth is determined by the bandwidth of the memory bus as well as the memory units.**
- **Memory bandwidth can be improved by increasing the size of memory blocks.**
 - **The underlying system takes l time units (where l is the latency of the system) to deliver b units of data (where b is the block size).**
- ◆ **Bandwidth has improved more than latency**
 - **23% per year vs 7% per year**

Impact of Memory Bandwidth: Example

- ◆ **In previous example, if block size is 4 words instead of 1 word.**
- ◆ **For the dot-product computation:**
 - **Assuming that the vectors are laid out linearly in memory, 8 FLOPs (4 multiply-adds) can be performed in 200 cycles.**
 - **A single memory access fetches 4 consecutive words in a vector.**
 - **2 accesses can fetch 4 elements of each of the vectors.**
 - **One FLOP every 25 ns, for a peak speed of 40 MFLOPS**
- ◆ **Bandwidth helps to hide (some) latency**

Impact of Memory Bandwidth

- ◆ **Note that increasing block size does not change latency of the system.**
- ◆ **Physically, it can be viewed as a wide data bus (4 words or 128 bits) connected to multiple memory banks.**
- ◆ **In practice, such wide buses are expensive to construct.**
- ◆ **In a more practical system, consecutive words are sent on the memory bus on subsequent bus cycles after the first word is retrieved.**

Impact of Memory Bandwidth

◆ In order to take advantage of full bandwidth

- The data layouts were assumed to be such that consecutive data words in memory were used by successive instructions
- Spatial locality of reference

◆ If spatial locality is poor

- Computations often need to be reordered to enhance spatial locality of reference.

Impact of Memory Bandwidth: Example

Consider the following code fragment:

```
for (i = 0; i < 1000; i++)  
    column_sum[i] = 0.0;  
    for (j = 0; j < 1000; j++)  
        column_sum[i] += b[j][i];
```

The code fragment sums columns of the matrix `b` into a vector `column_sum`.

Potential Problem

- ◆ The vector `column_sum` is small and easily fits into the cache
- ◆ The matrix `b` is accessed in a column order.
- ◆ What if the matrix is stored in row major (often the case)
 - The strided access results in very poor performance.

How to Fix the Problem?

We can fix the above code as follows:


```
for (i = 0; i < 1000; i++)  
    column_sum[i] = 0.0;  
for (j = 0; j < 1000; j++)  
    for (i = 0; i < 1000; i++)  
        column_sum[i] += b[j][i];
```

In this case, the matrix is traversed in a row-order and performance can be expected to be significantly better.


Memory Bandwidth Utilization: Another Example

◆ Matrix-vector multiplication

$$\mathbf{A} \mathbf{x} = \mathbf{b}$$



(a) Column major data access



(b) Row major data access.

Multiplying a matrix with a vector: (a) multiplying column-by-column, keeping a running sum; (b) computing each element of the result as a dot product of a row of the matrix with the vector.

Other Approaches for Hiding Memory Latency

- ◆ **Example: When browsing the web on a very slow network connection, one may:**
 - anticipate which pages we are going to browse ahead of time and issue requests for them in advance;
 - open multiple browsers and access different pages in each browser, thus while we are waiting for one page to load, we could be reading others; or
- ◆ **The first approach is called *prefetching*, the second *multithreading*.**

Multithreading for Latency Hiding

A thread is a single stream of control in the flow of a program.

We illustrate threads with a simple example:

```
for (i = 0; i < n; i++)  
    c[i] = dot_product(get_row(a, i), b);
```

Each dot-product is independent of the other, and therefore represents a concurrent unit of execution.

We can safely rewrite the above code segment as:

```
for (i = 0; i < n; i++)  
    c[i] = create_thread(dot_product, get_row(a, i), b);
```

Multithreading for Latency Hiding

- ◆ **The previous example assumes:**
 - **The memory system is capable of servicing multiple outstanding requests, and**
 - **The processor is capable of switching threads at every cycle.**
- ◆ **It also requires the program to have an explicit specification of concurrency in the form of threads.**
- ◆ **GPUs and multi-threaded CPUs can switch the context of execution in every cycle. Consequently, they are able to hide latency effectively.**

Prefetching for Latency Hiding

- ◆ **Misses on loads cause programs to stall.**
- ◆ **Why not advance the loads so that by the time the data is actually needed, it is already there!**
- ◆ **The only drawback is that you might need more space to store advanced loads.**
- ◆ **However, if the advanced loads are overwritten, we are no worse than before!**

Tradeoffs of Multithreading and Prefetching

- ◆ Multithreading and prefetching are critically impacted by the memory bandwidth.

- ◆ Assume

- A processor with 1 GHz clock, 32K cache with line size of 16 bytes, single cycle access to the cache, and 100 ns latency to DRAM.
- The computation has a cache hit ratio at
 - 1 KB of 25%; 32 KB of 90%.
- one data request in every 10 cycles

- ◆ 1-thread and 32-threads cases

- A 1-thread execution in which the entire cache is available
 - Required bandwidth = $(1-90\%) * 16B / (10 * 1ns) = 160 \text{ MB/s}$
- A 32-thread execution where each thread has a cache residency of 1 KB
 - Required bandwidth = $(1-25\%) * 16B / (10 * 1ns) = 1.2\text{GB/s}$

Tradeoffs of Multithreading and Prefetching

- ◆ **Bandwidth requirements of a multithreaded system may increase very significantly because of the smaller cache residency of each thread.**
- ◆ **Multithreaded systems become bandwidth bound instead of latency bound.**
- ◆ **Multithreading and prefetching only address the latency problem and may often exacerbate the bandwidth problem.**
- ◆ **Multithreading and prefetching also require significantly more hardware resources in the form of storage.**

Summary

◆ Convergence in parallel architecture

- **Initially: close coupling of programming model and architecture**
 - shared address space, message passing. data parallel
- **Now: separation and identification of dominant models/architectures**
 - **Programming models:**
 - ◆ shared address space
 - ◆ message passing
 - ◆ data parallel
 - **Architectures:**
 - ◆ small-scale shared memory
 - ◆ large-scale distributed memory
 - ◆ large-scale SMP cluster

Summary

◆ Memory performance is critical:

- Exploiting spatial and temporal locality in applications is critical for amortizing memory latency and increasing effective memory bandwidth.
- The ratio of the number of operations to number of memory accesses is a good indicator of anticipated tolerance to memory bandwidth.
- Memory layouts and organizing computation appropriately can make a significant impact on the spatial and temporal locality.

◆ Latency and bandwidth are two related but different concepts

- Memory reuse helps both latency and bandwidth
- Prefetching and multi-threading can help to deal with latency assuming sufficient bandwidth