



# **Parallel & Distributed Computing**

## **Programming with OpenMP**

Lecture 4, Spring 2014

Instructor: 罗国杰

[gluo@pku.edu.cn](mailto:gluo@pku.edu.cn)

# **Overall Outline**

---

## ◆ **Basics of OpenMP**

- APIs, directives, etc.

## ◆ **Correctness**

- Barriers
- Race conditions
- Mutual Exclusion
- Memory fence

## ◆ **Performance**

- Locality
- Loop scheduling
- Loop transformations

# **What is OpenMP?**

---

## ◆ An abbreviation for

- Short version
  - Open Multi-Processing
- Long version
  - Open specifications for Multi-Processing via collaborative work between interested parties from the hardware and software industry, government and academia

## ◆ OpenMP is an API for parallel programming

- First developed by the OpenMP Architecture Review Board (1997), now a standard
- Designed for shared-memory multiprocessors
- Set of compiler directives, library functions, and environment variables, but not a language
- Can be used with C, C++, or Fortran
- Based on fork/join model of threads

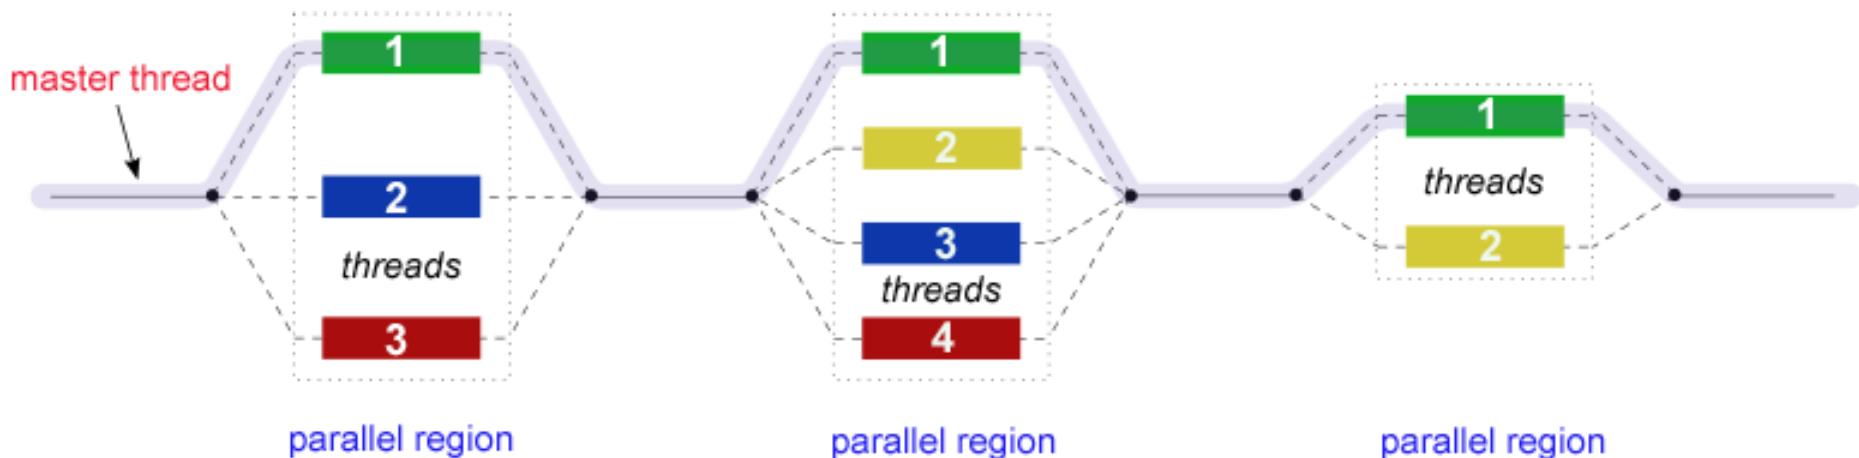
# Fork-Join Model

## ◆ FORK

- The master thread then **creates (or awakens)** a team of parallel *threads*.

## ◆ JOIN

- When the team threads complete the statements in the parallel region construct, they **synchronize and terminate**, leaving only the master thread.

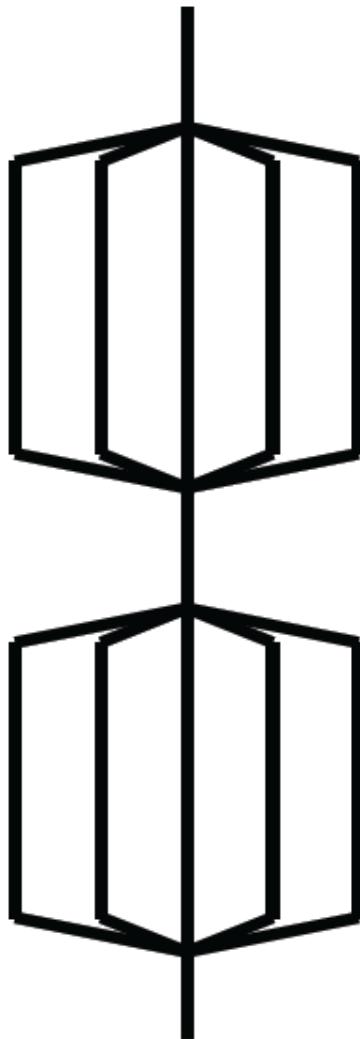


[http://en.wikipedia.org/wiki/Flynn's\\_taxonomy](http://en.wikipedia.org/wiki/Flynn's_taxonomy)

Blaise Barney (LLNL), Introduction to Parallel Computing

# *Relating Fork/Join to Code*

```
for {  
    [REDACTED]  
}  
[REDACTED]  
for {  
    [REDACTED]  
}  
[REDACTED]
```



Sequential code

Parallel code

Sequential code

Parallel code

Sequential code

# ***OpenMP Components***

---

## ◆ Three API components

- Compiler directives
- Runtime library routines
- Environment variables

[http://en.wikipedia.org/wiki/Flynn's\\_taxonomy](http://en.wikipedia.org/wiki/Flynn's_taxonomy)

Blaise Barney (LLNL), Introduction to Parallel Computing

# *Syntax of Compiler Directives*

---

## ◆ *pragma*: a C/C++ compiler directive

- (other compiler directives: #include, #define, ...)
- Stands for “pragmatic information”
- A way for the programmer to communicate with the compiler
- Pragmas are handled by the preprocessor
- Compilers are free to ignore pragmas

## ◆ All OpenMP pragmas have the syntax:

#pragma omp <directive-name> [clause, ...]

## ◆ Pragmas appear immediately *before* relevant construct

# Hello World

---

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[]) {
    int nthreads, tid;
    /* Fork a team of threads giving them their own copies of variables */
#pragma omp parallel private(nthreads, tid)
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Hello World from thread = %d\n", tid);
} /* All threads join master thread and disband */
}
```

# *Output – Non-deterministic!*

```
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
Hello World from thread = 5
Hello World from thread = 7
Number of threads = 16
Hello World from thread = 0
Hello World from thread = 8
Hello World from thread = 3
Hello World from thread = 1
Hello World from thread = 6
Hello World from thread = 15
Hello World from thread = 14
Hello World from thread = 10
Hello World from thread = 2
Hello World from thread = 9
Hello World from thread = 11
Hello World from thread = 13
Hello World from thread = 12
Hello World from thread = 4
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
Hello World from thread = 6
Hello World from thread = 10
Hello World from thread = 8
Number of threads = 16
Hello World from thread = 0
Hello World from thread = 5
Hello World from thread = 14
Hello World from thread = 15
Hello World from thread = 13
Hello World from thread = 12
Hello World from thread = 4
Hello World from thread = 3
Hello World from thread = 9
Hello World from thread = 2
Hello World from thread = 7
Hello World from thread = 1
Hello World from thread = 11
```

# ***Supplemental: printf & cout***

- ◆ C/C++ is not aware of “threads,” but POSIX is.
  - <http://pubs.opengroup.org/onlinepubs/9699919799/functions/flockfile.html>
  - All functions that reference (FILE \*) objects shall behave as if they use flockfile() and funlockfile() internally to obtain ownership of these (FILE \*) objects.

- ◆ What about replacing the printf statement ... ?

printf("Hello World from thread = %d\n", tid);

by

cout << "Hello World from thread = " << tid << endl;

# OpenMP Compilers Perform the Translations to Threads (e.g. pthreads)

```
int a, b;
main() {
    // serial segment
    #pragma omp parallel num_threads (8) private (a) shared (b)
    {
        // parallel segment
    }
    // rest of serial segment
}
```

Sample OpenMP program

```
int a, b;
main() {
    // serial segment
    for (i = 0; i < 8; i++)
        pthread_create (....., internal_thread_fn_name, ...);
    for (i = 0; i < 8; i++)
        pthread_join (.....);
    // rest of serial segment
}

void *internal_thread_fn_name (void *packaged_argument) {
    int a;
    // parallel segment
}
```

Code inserted by  
the OpenMP  
compiler

Corresponding Pthreads translation

# **Matching Threads with CPUs**

- ◆ Function *omp\_get\_num\_procs* returns the number of physical processors available to the parallel program

```
int omp_get_num_procs(void);
```

- ◆ Function *omp\_set\_num\_threads* allow you to set the number of threads that should be active in parallel sections of code

```
void omp_set_num_threads(int t);
```

- The function can be called with different arguments at different points in the program

# **Pragma: parallel for**

---

## ◆ The compiler directive

#pragma omp parallel for

**tells the compiler that the *for* loop which immediately follows can be executed in parallel**

- **The number of loop iterations must be computable at run time before loop executes**
- **Loop must not contain a *break*, *return*, or *exit***
- **Loop must not contain a *goto* to a label outside loop**

## ***Example: parallel for***

---

```
int a[1000], b[1000], s[1000];
```

...

```
#pragma omp parallel for
for (i = 0; i < 1000; i++)
    s[i] = a[i] + b[i];
```

- ◆ **Threads are assigned an independent set of iterations**
- ◆ **Threads must wait at the end of construct**

# Which Loop to Make Parallel?

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize a[][][], b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]  
    }  
}
```

# Which Loop to Make Parallel?

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++)          // Loop-carried dependences  
        for (i = 0; i < N; i++)      // Can execute in parallel  
            for (j = 0; j < N; j++)  // Can execute in parallel  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]
```

# *Minimizing Threading Overhead*

---

- ◆ There is a fork/join for every instance of

```
#pragma omp parallel for
```

```
for (...) {
```

```
...
```

```
}
```

- ◆ Since fork/join is a source of overhead, we want to maximize the amount of work done for each fork/join; i.e., the *grain size*
- ◆ Hence we choose to make the middle loop parallel

# *Almost Right, but Not Quite*

```
int main() {  
    int i, j, k;  
    float **a , **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for  
        for (i = 0; i < N; i++)  
            for(j = 0(j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]  
}
```

Problem: j is a shared variable

# ***Clause: private***

---

## ◆ **Clause**

- An optional, additional component to a pragma

## ◆ **Private clause: directs compiler to make one or more variables private**

```
#pragma omp ... private (<variable list>)
```

# **Problem Solved with private Clause**

```
int main() {  
    int i, j, k;  
    float **a , **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for private (j)  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]  
}
```

Tell compiler to make listed variables private

# *Another Example*

---

```
int i;  
float *a, *b, *c, tmp;  
...  
for (i = 0; i < N; i++) {  
    tmp = a[i] / b[i];  
    c[i] = tmp * tmp;  
}
```

- ◆ Loop is perfectly parallelizable except for shared variable *tmp*

# **Solution**

---

```
int i;  
float *a, *b, *c, tmp;  
...  
#pragma omp parallel for private (tmp)  
for (i = 0; i < N; i++) {  
    tmp = a[i] / b[i];  
    c[i] = tmp * tmp;  
}
```

# **More About Private Variables**

---

- ◆ Each thread has its own copy of the private variables
- ◆ If  $j$  is declared private, then inside the *for* loop no thread can access the “other”  $j$  (the  $j$  in shared memory)
- ◆ No thread can use a previously defined value of  $j$
- ◆ No thread can assign a new value to the shared  $j$
- ◆ Private variables are **undefined** at loop entry and loop exit, reducing execution time

# ***Clause: firstprivate***

---

- ◆ The *firstprivate* clause tells the compiler that the private variable should inherit the value of the shared variable upon loop entry
  - The value is assigned once per thread, not once per loop iteration

# **Example: firstprivate**

---

```
a[0] = 0.0;  
for (i = 1; i < N; i++)  
    a[i] = alpha(i, a[i-1]);
```

```
#pragma omp parallel for firstprivate (a)  
for (i = 0; i < N; i++)  
    b[i] = beta(i, a[i]);  
    a[i] = gamma(i);  
    c[i] = delta(a[i], b[i]);  
}
```

# **Clause: firstprivate**

---

## ◆ #pragma omp ... firstprivate(x)

- **$x$  is a fundamental data type**
  - Private  $x$  is directly copied from the shared  $x$
- **$x$  is an array**
  - Copy the data with  $\text{sizeof}(x)$  to the private memory
- **$x$  is a pointer**
  - Private  $x$  points to the same location as the shared  $x$
- **$x$  is a class instance**
  - Copy constructor is called to create the private  $x$

# ***Clause: lastprivate***

---

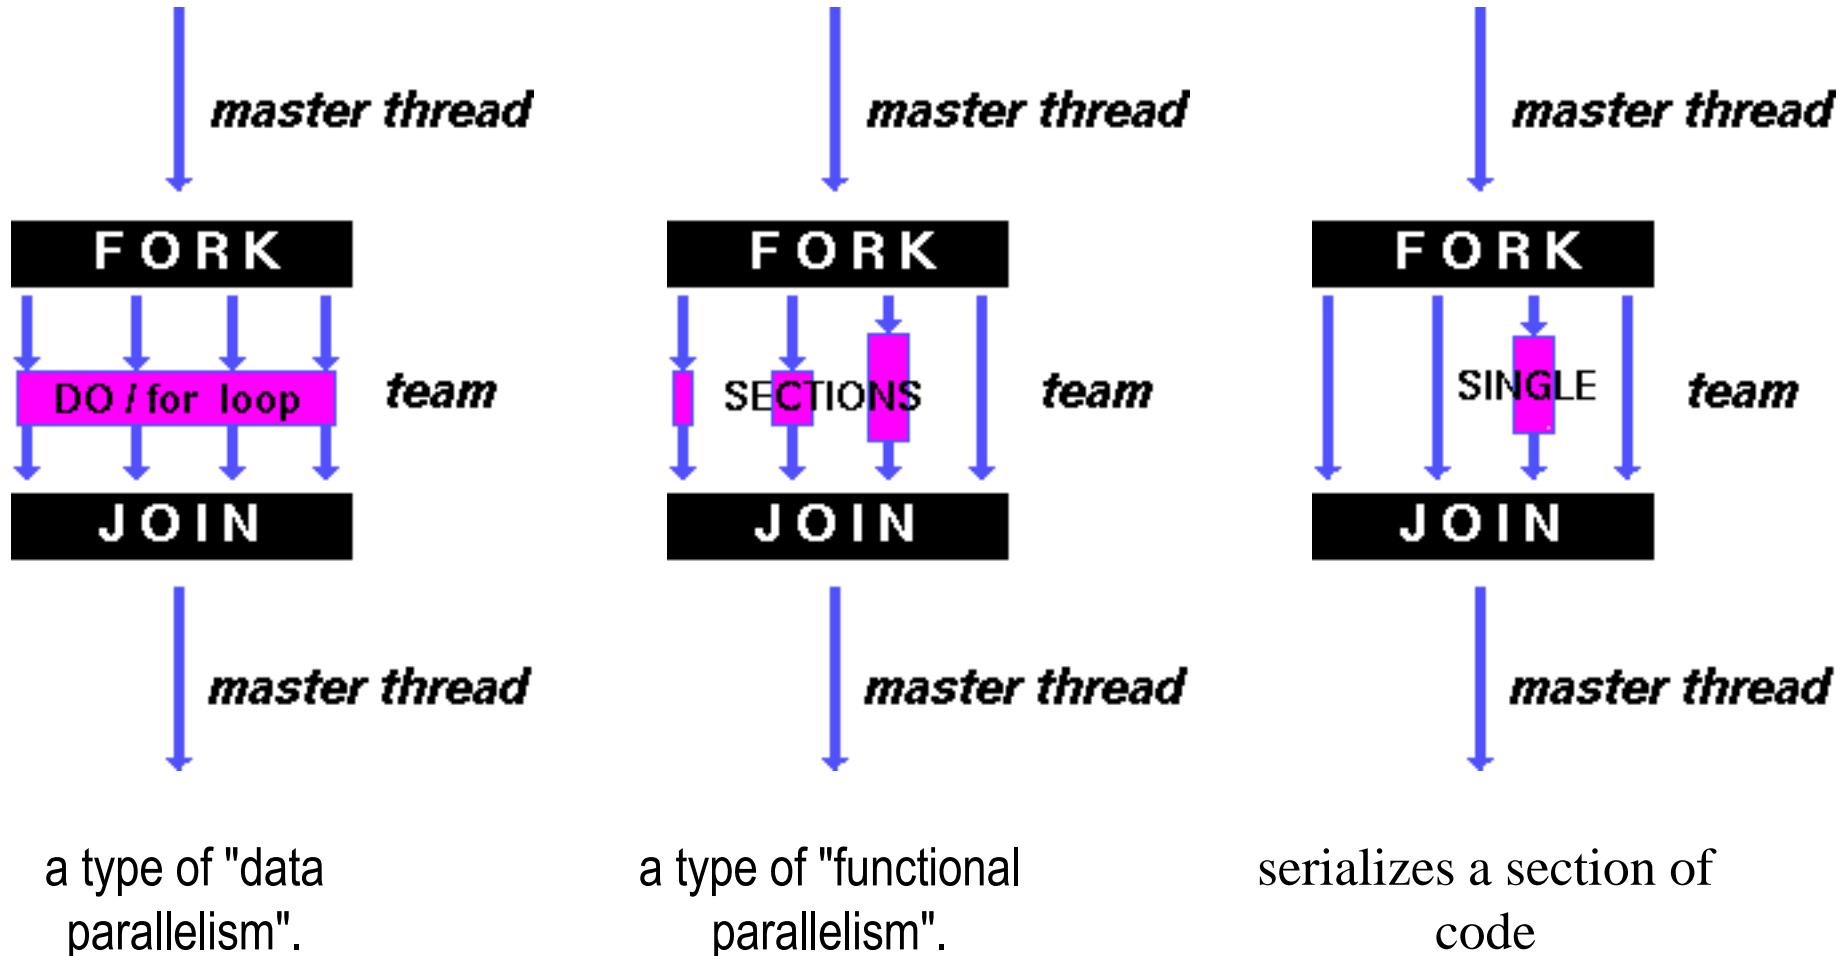
- ◆ The *lastprivate* clause tells the compiler that the value of the private variable after the *sequentially last* loop iteration should be assigned to the shared variable upon loop exit
  - In other words, when the thread responsible for the sequentially last loop iteration exits the loop, its copy of the private variable is copied back to the shared variable

# **Example: lastprivate**

---

```
#pragma omp parallel for lastprivate (x)
for (i = 0; i < N; i++)
    x = foo(i);
    y[i] = bar(i, x);
}
last_x = x; // == foo(N-1)
```

# Work-Sharing Constructs



# *Pragma: parallel*

---

- ♦ In the effort to increase grain size, sometimes the code that should be executed in parallel goes beyond a single *for* loop
  - The *parallel* pragma is used when a block of code should be executed in parallel
  - SPMD-style programming
    - Single program, multiple data

# **Pragma: for**

---

- ◆ The **for** pragma is used inside a block of code already marked with the *parallel* pragma
  - It indicates a *for* loop whose iterations should be divided among the active threads
  - There is a *barrier synchronization* of the threads at the end of the *for* loop

# **Pragma parallel and Pragma for**

## ◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads

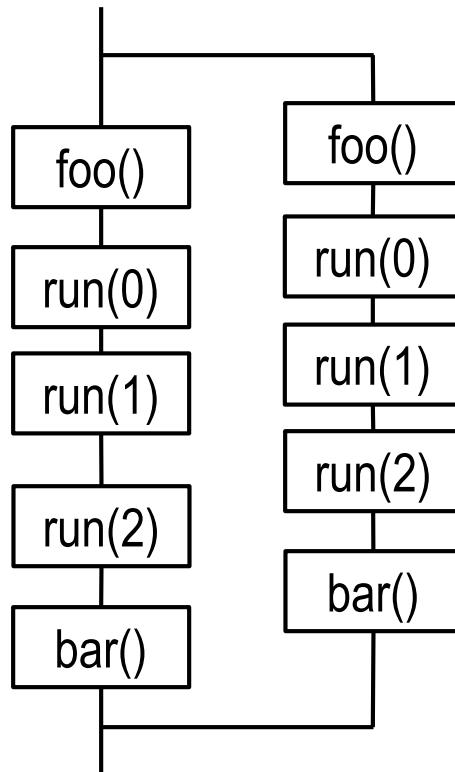
```
#pragma omp parallel \
num_threads(2)
{
    foo();
    #pragma omp for
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```

```
#pragma omp parallel \
num_threads(2)
{
    foo();
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```

# ***Pragma parallel and Pragma for***

## ◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads



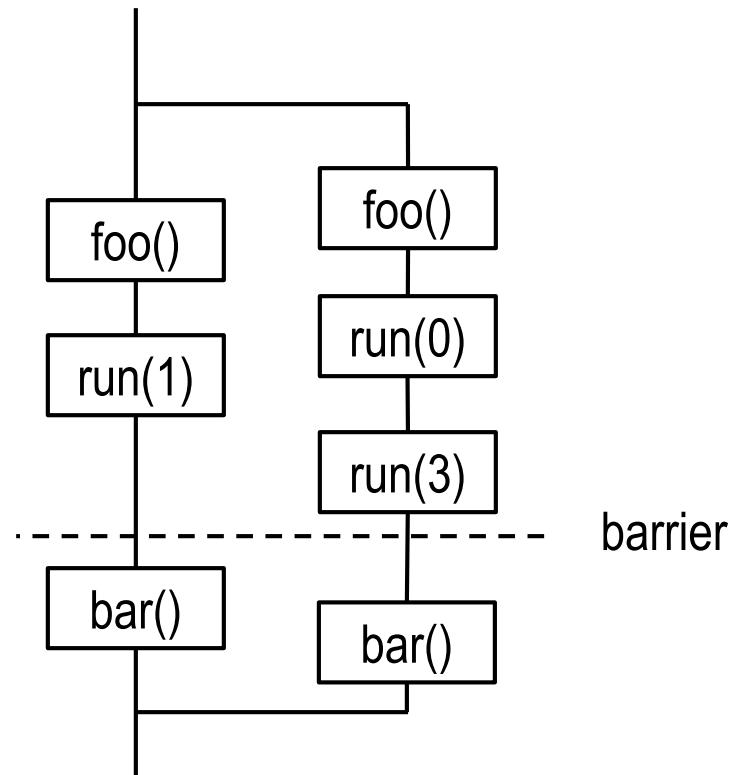
```
#pragma omp parallel \
num_threads(2)
{
    foo();
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```

# **Pragma parallel and Pragma for**

## ◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads

```
#pragma omp parallel \
num_threads(2)
{
    foo();
    #pragma omp for
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```



# *Parallelize For Loop*

---

# **Pragma: single**

---

- ◆ **The *single* pragma is used inside a parallel block of code**
  - It tells the compiler that only a single thread should execute the statement or block of code immediately following
  - May be useful when dealing with sections of code that are not thread safe (such as I/O)
  - Threads in the team that do not execute the single directive, wait at the end of the enclosed code block, unless a nowait clause is specified.

# **Example: master and single nowait**

```
tid = omp_get_thread_num();
if (tid == 0) {
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

=

```
#pragma omp master
{
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

≈

```
#pragma omp single nowait
{
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

# ***Pragma: sections and section***

---

- ◆ Directive sections specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- ◆ Independent section directives are nested within a sections directive.
  - Each section is executed once by a thread in the team.
  - Different sections may be executed by different threads.
  - It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.

# **Example: sections and section**

```
#pragma omp parallel shared(a,b,c,d) private(i)
{
# pragma omp sections
{
# pragma omp section
{
    for (i=0; i<N; i++)
        c[i] = a[i] + b[i];
}
# pragma omp section
{
    for (i=0; i<N; i++)
        d[i] = a[i] * b[i];
}
} /* end of sections */
} /* end of parallel section */
```

# **Clause: reduction**

---

- ◆ Reductions are so common that OpenMP provides a reduction clause for the *parallel*, *for*, and *sections*

```
#pragma omp ... reduction (op : list)
```

- A **private** copy of each list variable is created and initialized depending on the *op*
  - The identity value *op* (e.g., 0 for addition)
- These copies are **updated locally** by threads
- At end of construct, local copies are combined through *op* into a single value and combine the value in the original **shared** variable

# C/C++ Reduction Operation

- ◆ Reduction with an **associative** binary operator  $\oplus$

$$a_1 \oplus a_2 \oplus a_3 \oplus \dots \oplus a_n$$

- ◆ A range of associative and commutative operators can be used with reduction
- ◆ Initial values are the ones that make sense

Operator	Initial Value
+	0
*	1
-	0
$\wedge$	0

Operator	Initial Value
$\&$	$\sim 0$
$ $	0
$\&\&$	1
$\ $	0

# Reduction: an Artificial Example

```
int sum = 3;
int prod = 5;
#pragma omp parallel for \
reduction(+:sum) \
reduction(*:prod) \
num_threads(2)
for (int i=0; i < 3; ++i) {
    int tid =
        omp_get_thread_num();
    sum += i;
    prod += i;
    printf("thread(%d) "
           "sum=%d prod=%d\n",
           tid, sum, prod);
}
printf("results: "
       "sum=%d prod=%d\n",
       tid, sum, prod);
```

## ◆ Assume

- **thread 0 executes the 1st and 2nd iterations, and**
- **thread 1 executes the 3<sup>rd</sup> iteration**

## ◆ Possible outputs

thread(0) sum=0 prod=1  
thread(1) sum=2 prod=3  
thread(0) sum=1 prod=2  
results: sum=6 prod=30

# Reduction: an Artificial Example

```
int sum = 3;
int prod = 5;
#pragma omp parallel for \
reduction(+:sum) \
reduction(*:prod) \
num_threads(2)
for (int i=0; i < 3; ++i) {
    int tid =
        omp_get_thread_num();
    sum += i;
    prod *= i;
    printf("thread(%d) "
           "sum=%d prod=%d\n",
           tid, sum, prod);
}
printf("results: "
       "sum=%d prod=%d\n",
       tid, sum, prod);
```

- ◆ **initial  $\text{sum}_{\text{private}} = 0$** 
  - **for the reduction of +**
- ◆ **initial  $\text{prod}_{\text{private}} = 1$** 
  - **for the reduction of \***
- ◆ **At the end of “parallel for” with reduction**
  - $\text{sum}_{\text{shared}} += \sum_{\text{thread}} \text{sum}_{\text{private}}$
  - $\text{prod}_{\text{shared}} *= \prod_{\text{thread}} \text{prod}_{\text{private}}$

# ***Strengths and Weaknesses of OpenMP***

---

## ◆ **Strengths**

- **Incremental parallelization & sequential equivalence**
- **Well-suited for domain decompositions**
- **Available on \*nix and Windows**

## ◆ **Weaknesses**

- **Not well-tailored for functional decompositions**
- **Compilers do not have to check for such errors as deadlocks and race conditions**